



Actual parameter values and their meaning | channel_neural

These axes identify features, coefficients, hidden units, or LL/LH/HH responses—not spatial coordinates.
Physical maps come from transport and signed products; the learned/fitted mixer supplies only three gains.




Actual response mapping

LL: logit -0.00522957
gain = 0.99607784

LH: logit -0.00522041
gain = 0.99608475

HH: logit +0.00522989
gain = 1.00392234

$g = 1 + 0.75 \cdot \tanh(\text{logit})$
Admissible gain: (0.25, 1.75)
Bounded gain ≠ stability proof.



Do the gains visibly change this output?

FP32 endpoint max change: 0.000000000e+00

Unit-gain RMS error: 6.544207081e-08
Selected RMS error: 6.544207081e-08

Small gain changes may fall below FP32 endpoint rounding. A rescaled channel map does not establish a visible or useful complete-solution improvement.

Lower error is better; this is one field.