



Stage compute and frozen budgets




Complete experiment verdicts





A required cost failure can make an otherwise accurate experiment BAD. Scores measure evidence attainment, not model quality.

Correctness math gap and utility checks




Training loss by update




Validation loss by update



track="continuum"





track="discrete"




Validation error versus training time




Validation error versus examples seen



Model parameter counts




Baseline adequacy and checkpoint selections




Initialization selection is not evidence of learned improvement. A weak trained control does not establish superiority over a published method.

Data efficiency validation observations











All 26 finite points; unavailable/nonpositive 4




All 26 finite points; unavailable/nonpositive 4


Work precision rms error with reference uncertainty




Work precision max error with reference uncertainty




All 72 finite points; unavailable/nonpositive 0




All 72 finite points; unavailable/nonpositive 0




All 52 finite points; unavailable/nonpositive 0




All 52 finite points; unavailable/nonpositive 0




All 72 finite points; unavailable/nonpositive 0



All 72 finite points; unavailable/nonpositive 0




All 52 finite points; unavailable/nonpositive 0






All 52 finite points; unavailable/nonpositive 0

Stress regime and tolerance coverage



Normal, best and worst cases use frozen field regimes, not retrospective exclusion. Cells with no target evidence remain NA.


Mean error across spatial resolution



All 144 finite points; unavailable/nonpositive 0




All 104 finite points; unavailable/nonpositive 0

All 144 finite points; unavailable/nonpositive 0



All 104 finite points; unavailable/nonpositive 0


Centered error across spatial resolution



All 144 finite points; unavailable/nonpositive 0


All 104 finite points; unavailable/nonpositive 0

All 144 finite points; unavailable/nonpositive 0




All 104 finite points; unavailable/nonpositive 0


High frequency error across spatial resolution




All 144 finite points; unavailable/nonpositive 0



All 104 finite points; unavailable/nonpositive 0




All 144 finite points; unavailable/nonpositive 0



All 104 finite points; unavailable/nonpositive 0


Scaling latency and throughput



All 10 finite points; unavailable/nonpositive 0


All 10 finite points; unavailable/nonpositive 0

All 10 finite points; unavailable/nonpositive 0




All 10 finite points; unavailable/nonpositive 0

Scaling peak GPU memory



Accuracy is the worst accepted-reference field in each distinct batch; missing references are NA. GPU memory is a shared paired-workload peak, not attributable to a family or shared host RAM.


Scaling accepted reference accuracy



All 10 finite points; unavailable/nonpositive 0

All 10 finite points; unavailable/nonpositive 0

All 10 finite points; unavailable/nonpositive 0



All 10 finite points; unavailable/nonpositive 0

Deployment complete cost components

track=NA | final_time=NA | mode=NA

NA — no verified measurements for this panel


First three components sum to attributed_seconds for a representative run; cost_seconds is independent complete warmed median. Do not sum these totals together. Missing or blocked deployment remains NA.

Policy coverage and false acceptance

track=NA | final_time=NA | mode=NA


NA — no verified measurements for this panel

Reference uncertainty and unresolved evidence




A tighter solver error beneath reference uncertainty is unresolved. Spatial and time refinement are not interchangeable.


First invocation versus randomized warmed timing




All 144 finite points; unavailable/nonpositive 0



All 104 finite points; unavailable/nonpositive 0




All 144 finite points; unavailable/nonpositive 0




All 104 finite points; unavailable/nonpositive 0


Independent field intervals error_ratio



Independent field intervals speed_ratio



Five gate evidence and missing coverage



Teal GOOD, coral BAD, gray NA. Scores are predeclared check-weight attainment, never scientific quality, effect size, or a proof.

Independent function risk and available evidence

track=NA | mode=NA

NA — no verified measurements for this panel

Supported amortization including offline costs

track=NA

Unestablished:

NA — no verified measurements for this panel

Every native software test execution

NA — no native per-stage JUnit artifacts were supplied