



Stage compute and frozen budgets



Each physical allocation appears once; batch/step records are never added. Sunk/origin costs are separate from current work. Failed stage costs remain visible; running report allocation and tariffs remain NA.


Physical confirmation coverage



Unmerged partial confirmation diagnostics


NA: no unmerged sealed partial endpoint errors

Complete experiment verdicts




A required cost failure can make an otherwise accurate experiment BAD. Scores measure evidence attainment, not model quality.

Correctness math gap and utility checks





Every individual check




161846 cells; no downsampling. Teal GOOD, coral BAD, gray NA, white padding. Exact identities and values in [check-cell-index.json.gz](#).

Training loss by update




Validation loss by update



track="continuum"





track="discrete"




Validation error versus training time



Validation error versus examples seen




Model parameter counts




Identical parameter counts collapse only for display by family/track; every model ID is retained in chart data. Parameter count is not FLOPs or runtime.

Baseline adequacy and checkpoint selections




Initialization selection is not evidence of learned improvement. A weak trained control does not establish superiority over a published method.

Data efficiency validation observations











All 176 finite points; unavailable/nonpositive 6274















All 176 finite points; unavailable/nonpositive 6274

Work precision rms error with reference uncertainty







Work precision max error with reference uncertainty









paired candidate / control











paired candidate / control




paired candidate / control


Stress regime and tolerance coverage




Mean error across spatial resolution




All 6336 finite points; unavailable/nonpositive 0



All 4608 finite points; unavailable/nonpositive 0




All 6336 finite points; unavailable/nonpositive 0




All 4608 finite points; unavailable/nonpositive 0


Centered error across spatial resolution




All 6336 finite points; unavailable/nonpositive 0



All 4608 finite points; unavailable/nonpositive 0




All 6336 finite points; unavailable/nonpositive 0




All 4608 finite points; unavailable/nonpositive 0


High frequency error across spatial resolution




All 6336 finite points; unavailable/nonpositive 0



All 4608 finite points; unavailable/nonpositive 0




All 6336 finite points; unavailable/nonpositive 0




All 4608 finite points; unavailable/nonpositive 0


Scaling latency and throughput




All 15 finite points; unavailable/nonpositive 0



All 15 finite points; unavailable/nonpositive 0




All 15 finite points; unavailable/nonpositive 0




All 15 finite points; unavailable/nonpositive 0


Scaling peak GPU memory




Scaling accepted reference accuracy




All 15 finite points; unavailable/nonpositive 0



All 15 finite points; unavailable/nonpositive 0



All 15 finite points; unavailable/nonpositive 0



All 15 finite points; unavailable/nonpositive 0

Deployment complete cost components

track=NA | final_time=NA | mode=NA

NA — no verified measurements for this panel


First three components sum to attributed_seconds for a representative run; cost_seconds is independent complete warmed median. Do not sum these totals together. Missing or blocked deployment remains NA.

Policy coverage and false acceptance

track=NA | final_time=NA | mode=NA


NA — no verified measurements for this panel

Reference uncertainty and unresolved evidence




A tighter solver error beneath reference uncertainty is unresolved. Spatial and time refinement are not interchangeable.


First invocation versus randomized warmed timing




All 6336 finite points; unavailable/nonpositive 0



All 4608 finite points; unavailable/nonpositive 0




All 6336 finite points; unavailable/nonpositive 0




All 4608 finite points; unavailable/nonpositive 0


Independent field intervals error_ratio



Independent field intervals speed_ratio



Five gate evidence and missing coverage



Teal GOOD, coral BAD, gray NA. Scores are predeclared check-weight attainment, never scientific quality, effect size, or a proof.

Independent function risk and available evidence

track=NA | mode=NA

NA — no verified measurements for this panel

Supported amortization including offline costs

track=NA

Unestablished:

NA — no verified measurements for this panel

Every native software test execution

